DEPTH

A drone has to know how far.

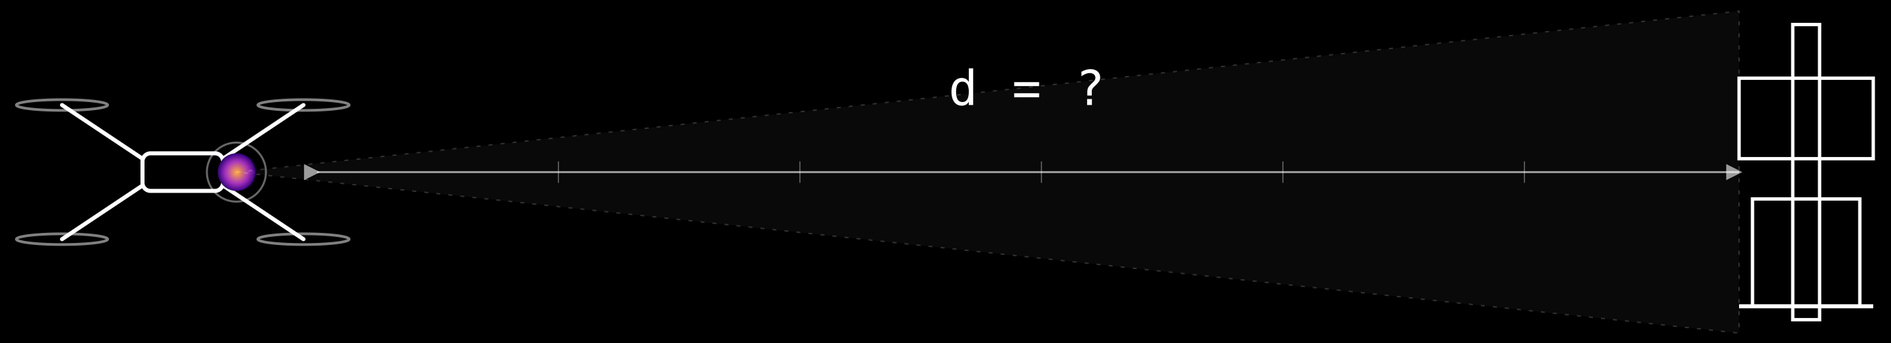

CAMERA → TARGET — EVERY LOOP DEPENDS ON THIS NUMBER

THE CURRENT TOOLBOX

Every option has a tax.

Stereo

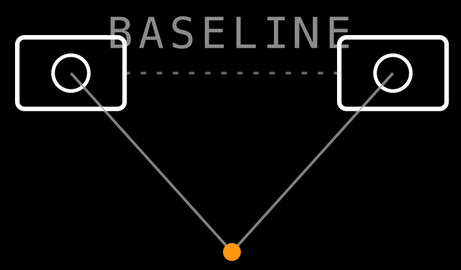

TAX

Two cameras. Calibration drift. Range collapses with baseline.

LiDAR

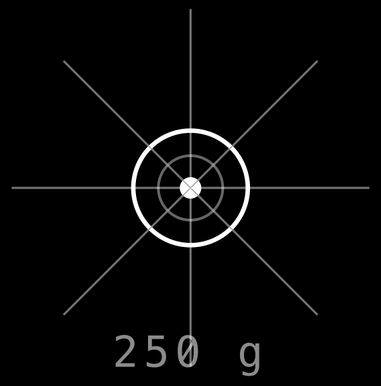

TAX

Heavy, expensive. Active emitter. Unfit for single-use.

Laser

TAX

One point only. Active emission — a beacon for jammers.

CLOSE ONE EYE.

You still know how far the wall is.

ONE EYE

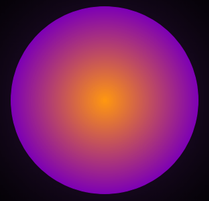

DEPTH, RECOVERED

CYCLOPS VISION

ONE EYE. ALL THE DEPTH YOU NEED.

A software-defined depth sensor.

ONE CAMERA. ALL THE DEPTH.

No baseline. No spinning mirrors. No emission.
Just a model — and the smallest board we could find.

ENCODER–DECODER
FOUNDATION BACKBONE
FINE–TUNED FOR UAV
DEPLOYED ON PI 5

● LIVE ON THE BENCH

Raspberry Pi 5

BATTERY–POWERED · NO GPU · NO ACCELERATOR

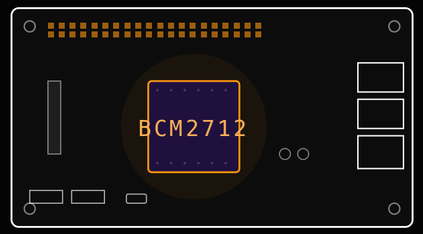

INFERENCE	5 fps live
QUANTIZATION	none yet
POWER	battery class
STATUS	running today

Standard vs finetuned output.

PORNHUB MODEL OPTIMIZATION

REFERENCE: DEPTH ANYTHING

RT-MONODEPTH-S: EVOLUTION

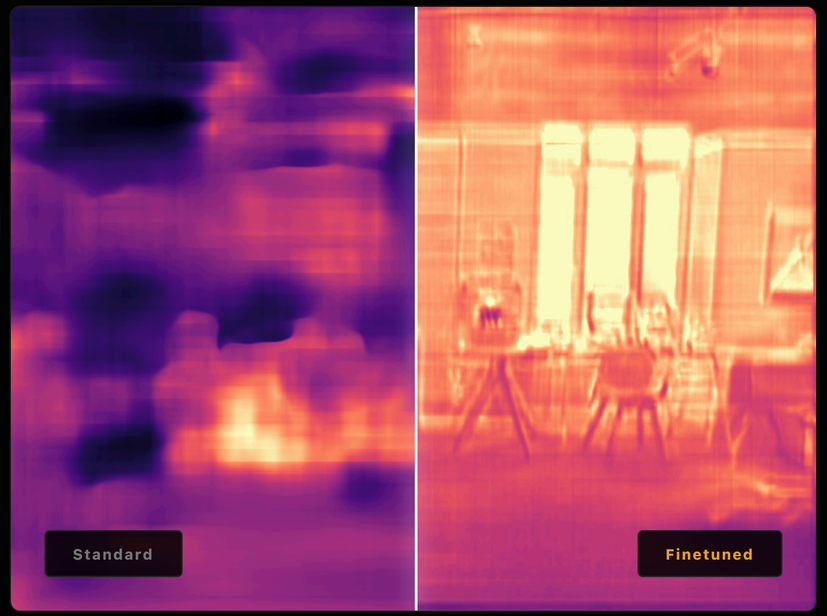

5 MB

is all you need to run state-of-the-art monocular depth — on the edge.

SOTA models ship at hundreds of megabytes and demand a GPU. We compressed the same task into a footprint that fits, runs, and survives on a Pi-class board.

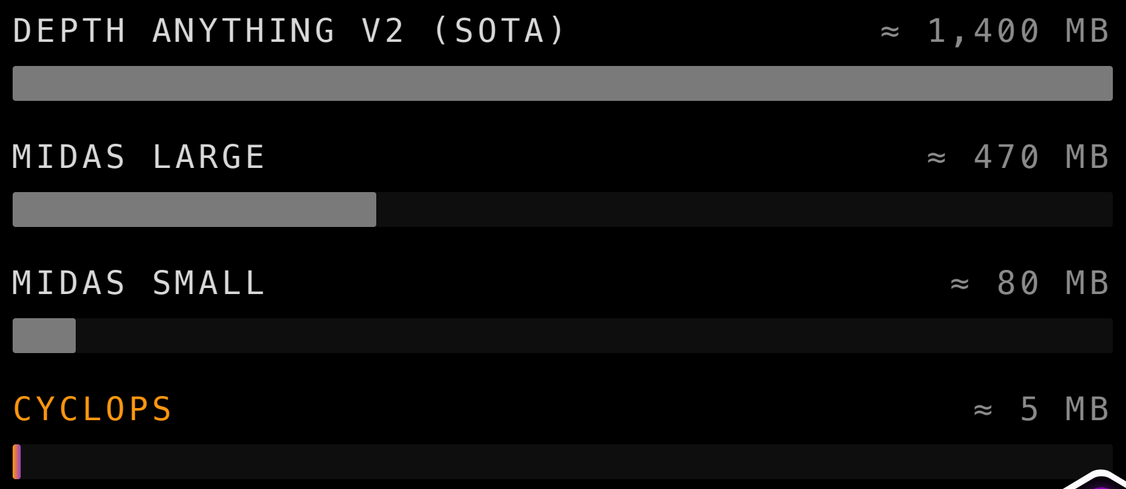

Three curves just crossed.

Loitering munitions go mainstream.

Attritable. Cost & weight rule.

The spectrum is being jammed.

Active emitters lose. Passive wins.

Foundation models fit on a Pi.

Edge inference is finally real.

Stereo accuracy. LiDAR range. Neither price.

STEREO

€400

- Two cameras + calibration
- Short effective range
- Fragile to vibration

LIDAR

€1,500+

- 250 g — too heavy
- 100 m range
- Active emitter
- Unfit for single-use

CYCLOPS VISION

camera + Pi

- Stereo-class accuracy
- LiDAR-class range
- Passive — no emission
- Single-board, attritable

Built in 48 hours by three.

Niclas

CYCLOPS VISION

Luca

CYCLOPS VISION

Lars

CYCLOPS VISION

CYCLOPS VISION

ONE EYE. ALL THE DEPTH YOU NEED.